



**NANYANG
TECHNOLOGICAL
UNIVERSITY**
SINGAPORE

NTU Academy for Professional
and Continuing Education

Ethical AI & Challenges

When you work with AI, the ethical call is still yours

Module 5

Created by :

Swarup Biswas
July 2026



"A computer can never be held
accountable."

"Therefore a computer must never make a management decision."

Source: IBM internal training manual, 1979

The accountability lands on you

AI amplifies a decision. It never owns it. The duty to check for harm stays with the human closest to the decision.

Companies trust AI. People don't.

Most organisations now use AI in real decisions — but few people trust it to be **fair and unbiased**.

That gap is where your judgment does its work.

Ethical harm enters through four doors

- **Bias** — is it repeating past discrimination, at scale?
- **Opacity** — can you explain the decision to the person it affects?
- **Privacy** — do you know what the tool does with people's data?
- **Misuse** — is the AI trusted beyond what it can actually do?

Bias: trained on our past, repeated at scale

Amazon — a hiring AI learned to penalise the word "women's" on a CV. Scrapped 2018.

Workday — a live 2026 lawsuit: AI screening accused of bias by race, age, disability.

HireVue — a deaf applicant rejected by an AI video interview.

Source: Mobley v. Workday (2025–26); ACLU/EEOC (2025)

The proxy trap: it measured spending, not sickness

A US health AI used **cost** as a stand-in for **need**. Black patients had less spent on them at the same illness level — so it under-flagged them.

*Patients flagged for extra care: **17.7%** Black. Fix the proxy → **46.5%**.*

Source: Obermeyer et al., Science, 2019

Opacity: "the algorithm decided" is not an explanation

When someone asks "why was I rated low?" you owe them a real answer.

Most AI tools cannot explain their own logic. A right answer for the wrong reason is still wrong.

Privacy: know what the tool does with the data

Before you feed a tool personal data, know: what it collects, who sees it, how long it's kept, and whether consent was real.

*Data pasted into a public AI tool may be kept, used for training, and **cannot be taken back**.*

Misuse: confident, fluent, and sometimes just wrong

Generative AI writes fluent, confident content that can be flatly wrong — invented facts, fake citations.

47% of enterprise AI users made a major decision on hallucinated content.

Source: 2024 enterprise survey

We over-trust AI — especially when we see it first

Automation bias: accepting the AI's answer even when your own knowledge says it's wrong.

*Read the AI's view first and you anchor to it. **Form your own view first.***

Audit a real AI — then defend its decision

- **A —** Ask any AI a real people-task: rank candidates, or write a rejection, or a review (use made-up names)
- Audit the output against the four doors: bias? can you explain it? privacy? made-up facts?
- **B —** In pairs, role-play: one is the person affected asking "why?"; the other must give a real answer — not "the system said so"
- Debrief: what made an answer honest instead of an abdication?

Five questions for any AI decision

1. Who does this affect — and did they know?
2. What is it really measuring?
3. Can I explain this to the person it affects?
4. Did I form my own view first?
5. If it goes wrong, how do we fix it?

What you do differently on Monday

Form your own view first — two minutes, before you read the AI's answer

Keep a plain decision note — what AI said, what you decided, why any difference

Three more habits that need no permission

Never deliver an outcome as "the system said so" — build the human explanation first

Verify before it reaches a person — AI output is a first draft

Tell people their right to a human review — GDPR and the EU AI Act give them one

The law is catching up

Framework	What it means for you
GDPR Art. 22	Right to a human review of automated decisions, and a real explanation
EU AI Act	Hiring & performance AI = high-risk; up to €15M / 3% turnover
US (EEOC)	Anti-discrimination law applies fully — "the algorithm did it" is no defence

Ethical AI isn't a policy to read — it's a call you make

*Every time an AI output reaches your desk, the ethics of what happens next are **yours**.*

Version Management

Month/Year	Version	Changes
May 2026	0.0	New Module Design
July 2026	1.0	Updated and Refreshed Content